

# Simple Cloud Cover Classification CNN

## Seminar Pattern Recognition

**Tishana Suthenthiran**

03.12.2025, University of Bern



# Introduction

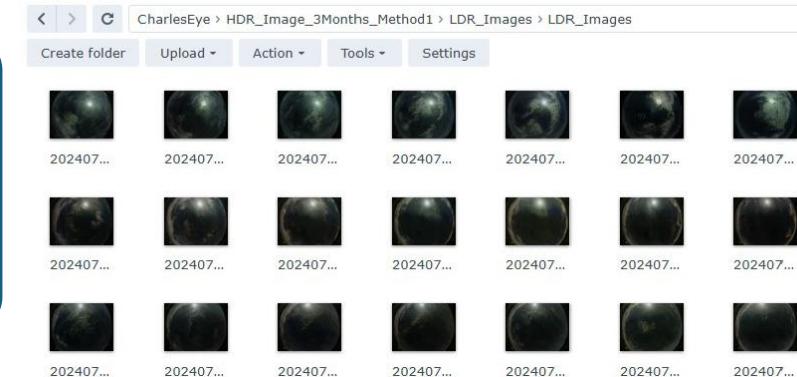
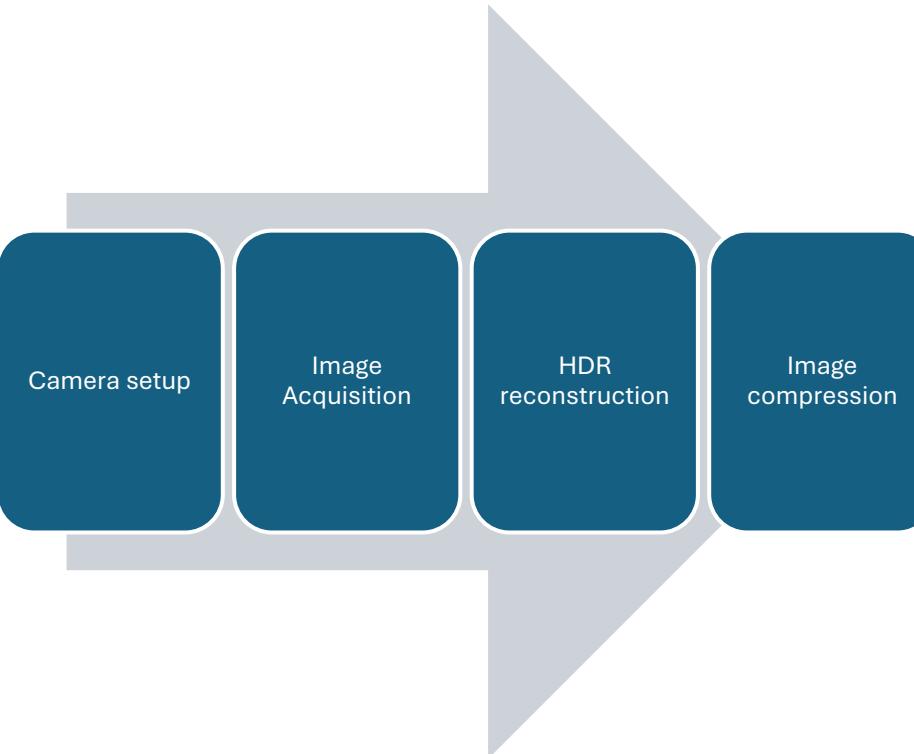
# GOALS OF THE PROJECT

- The aim of this project is to build a simple and effective **image classifier using convolutional neural networks (CNNs)**.
  - Goal 1: Work with an existing dataset of sky images.
  - Goal 2: Train a CNN for multi-class **classification of sky images**.
  - Goal 3: Evaluate the model's accuracy and present results .

The CNN will be used for further work in **irradiance prediction**.

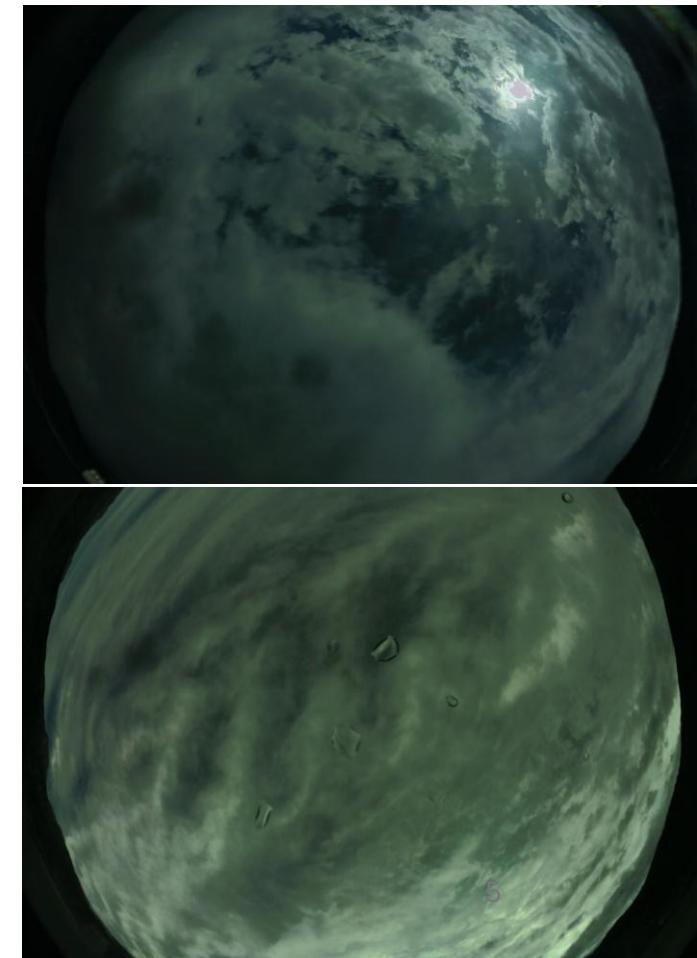
# Behind the Scenes

Fereshteh's PhD Project



# My Task: Cloud Cover Classification

- Images from July to end of September 2024 from 6am to 7.55pm
  - 12GB!!
  - We have a large dataset of sky images, but no automatic method to classify them into labels like
    - Clear
    - Cloudy
    - Overcast
- = CLOUD COVER

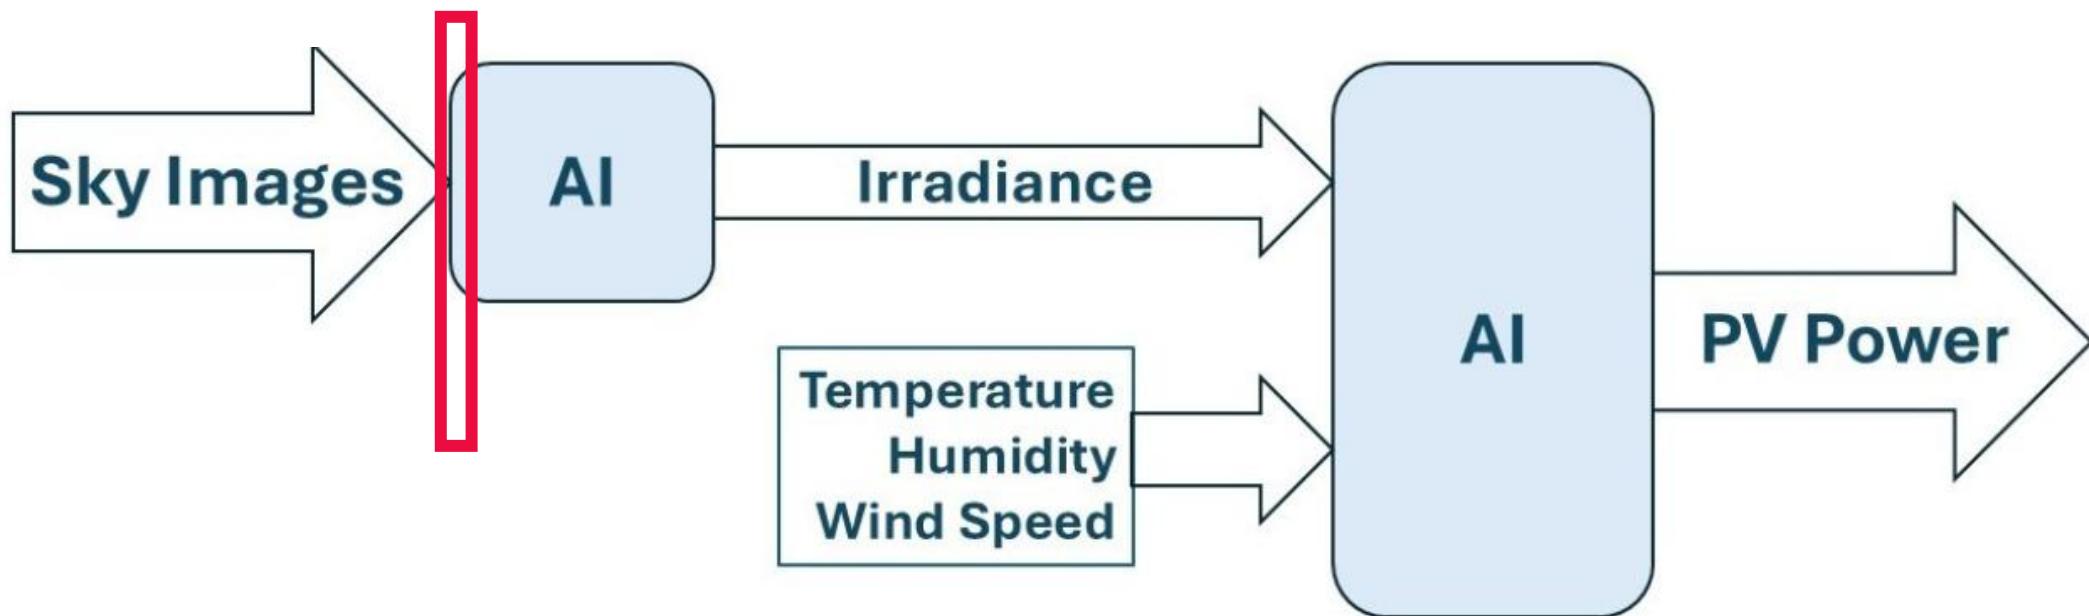


# Cloud cover $\leftrightarrow$ Irradiance

- The main source of uncertainty is **cloud movement**.
- Clouds dominate short-term PV variability and irradiance.
  - Classifying the sky gives immediate information about expected irradiance behaviour
  - Shadows reduce PV output by 50%-80% within minutes.
- Cloud classes describe distinct PV output patterns
  - Clear sky  $\rightarrow$  smooth, high irradiance
  - Cloudy sky  $\rightarrow$  rapid fluctuations, shading events
  - Overcast sky  $\rightarrow$  stable but low irradiance

# CNN as Preprocessing Step

- The CNN is not the final forecaster — it is the tool that turns raw images into useful inputs for downstream irradiance prediction.



# Preprocessing

# Preprocessing

- **Purpose**

- Ensure only clean, informative images enter the training dataset
- Remove frames that would confuse or degrade the CNN – sensitive to noise



$u^b$ 

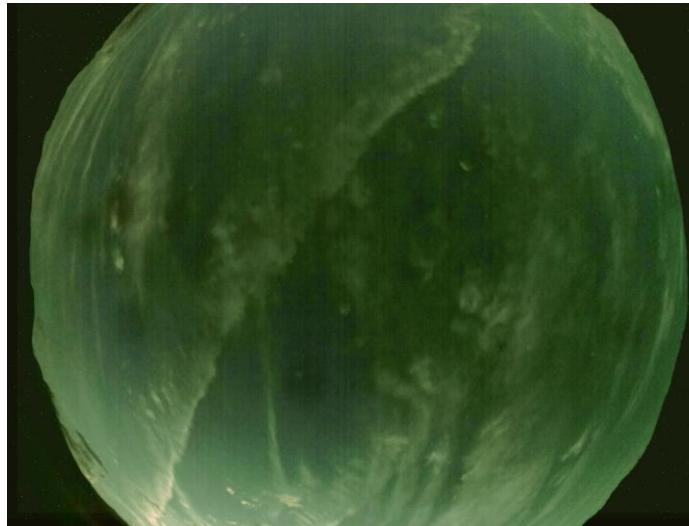
Time  
Window  
0600 - 0620

Early morning images are often

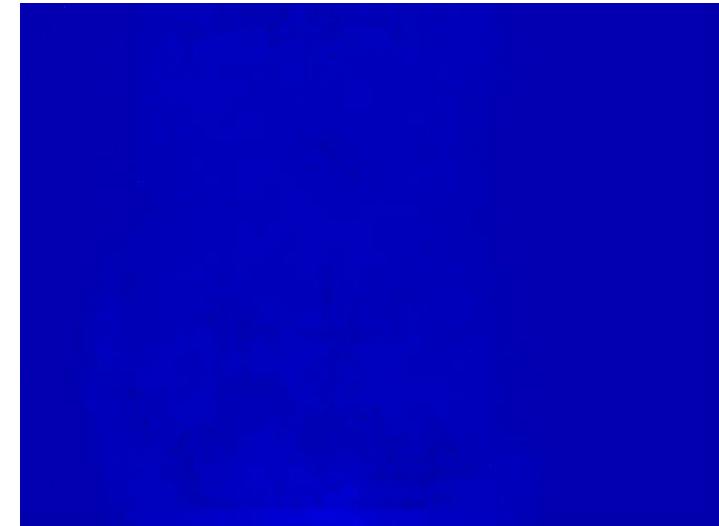
- Heavily color-shifted (purple/blue)
- Too dark
- Clouds not visible

→ Can confuse CNN during training

06.07.2024 06:00



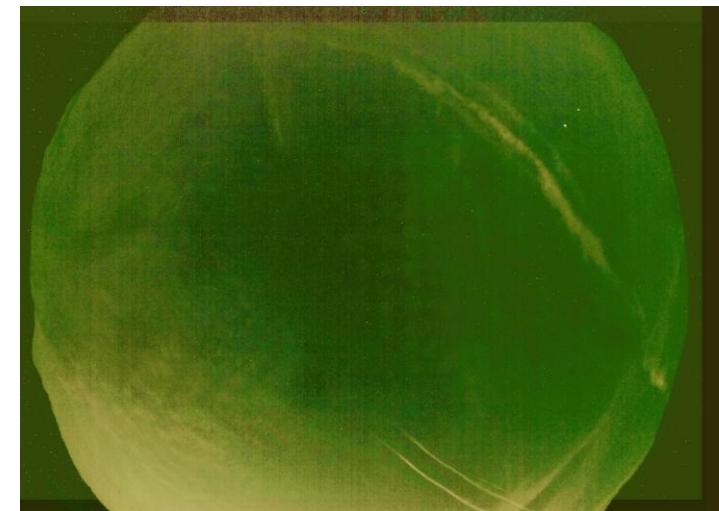
13.07.2024 06:20



06.07.2024 06:20

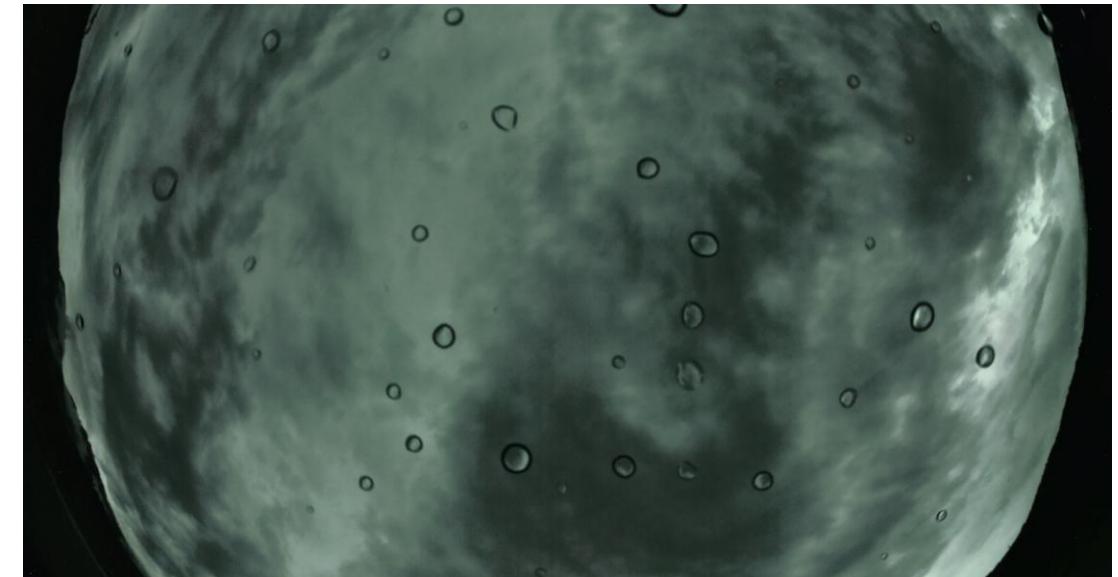
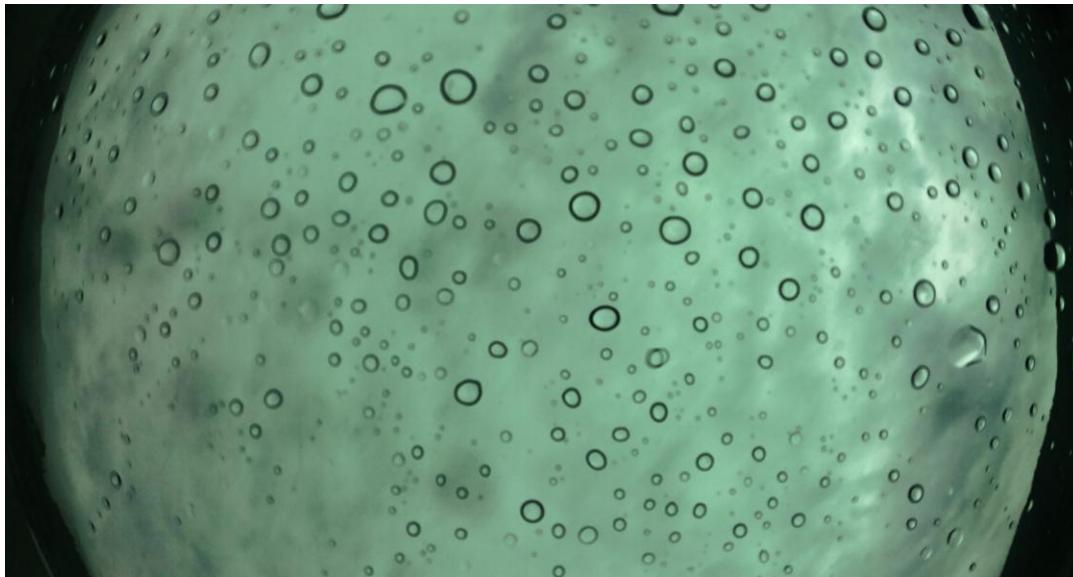


25.07.2024 06:10



Raindrop  
CNN + SVM

- Raindrops on the lens:
  - Distort the cloud structure
  - Create blurry spots unrelated to the sky
  - Can be mistaken as clouds
- Pretrained CNN + SVM classifier





## Color Dominance

### Algorithm:

1. Compute mean color value of each channel:
  - Blue, Green, Red
2. Compute each channel's ratio
  - $\text{ratios} = [\text{mean\_B}/\text{total}, \text{mean\_G}/\text{total}, \text{mean\_R}/\text{total}]$
3. Sort the ratios
  - $\text{top\_two} = \text{highest} + \text{second\_highest}$
4. If the top two channels take > 90% of the total colour:  
The image is **colour-dominated**  
→ move to quarantine/dark\_or\_tinted

02.09.2024 19:50

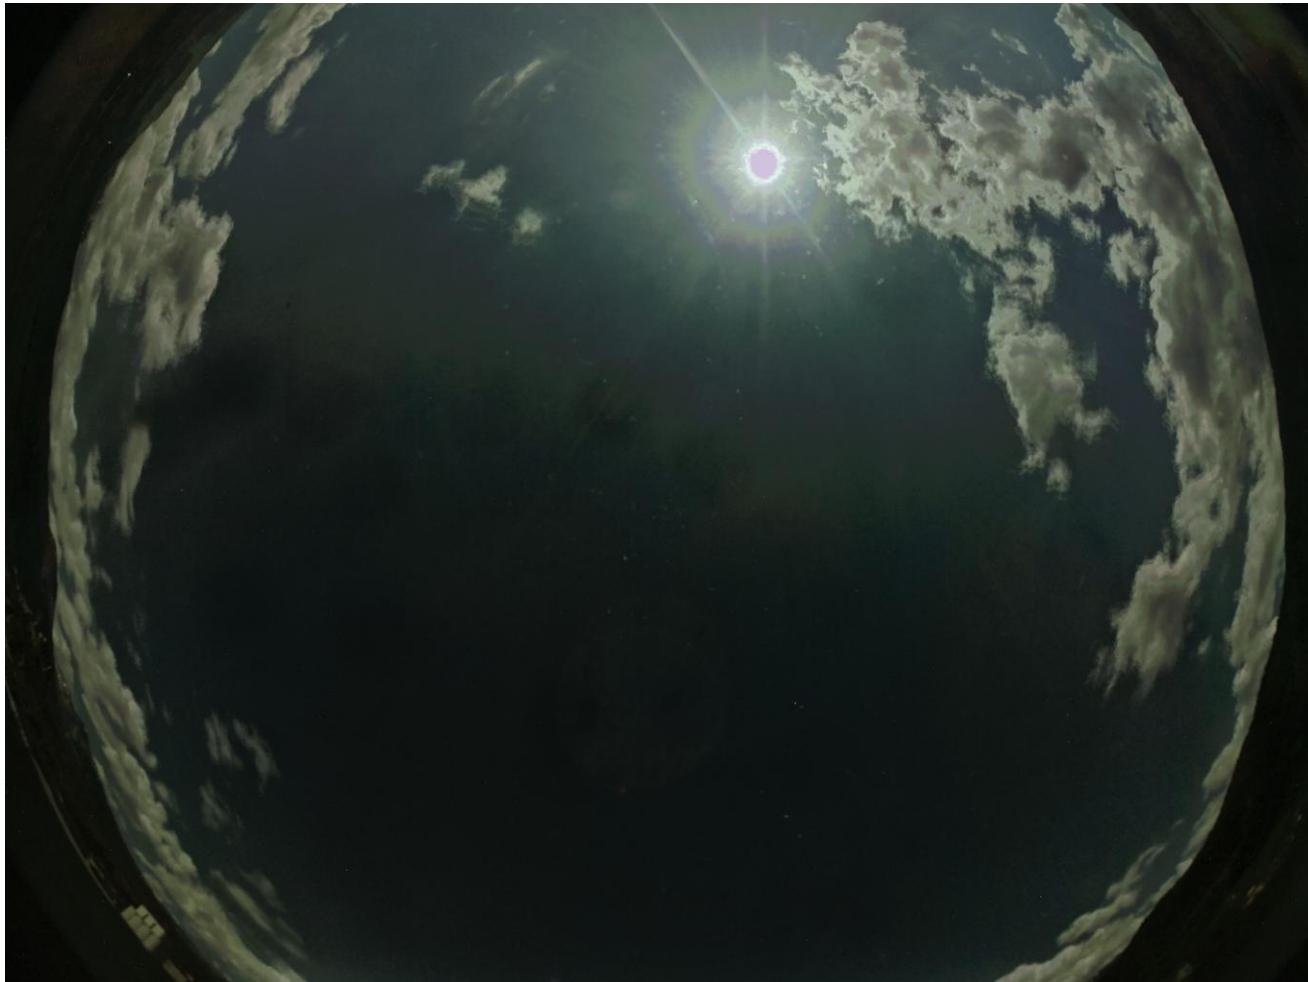
10.08.2024 06:40

# Labelling

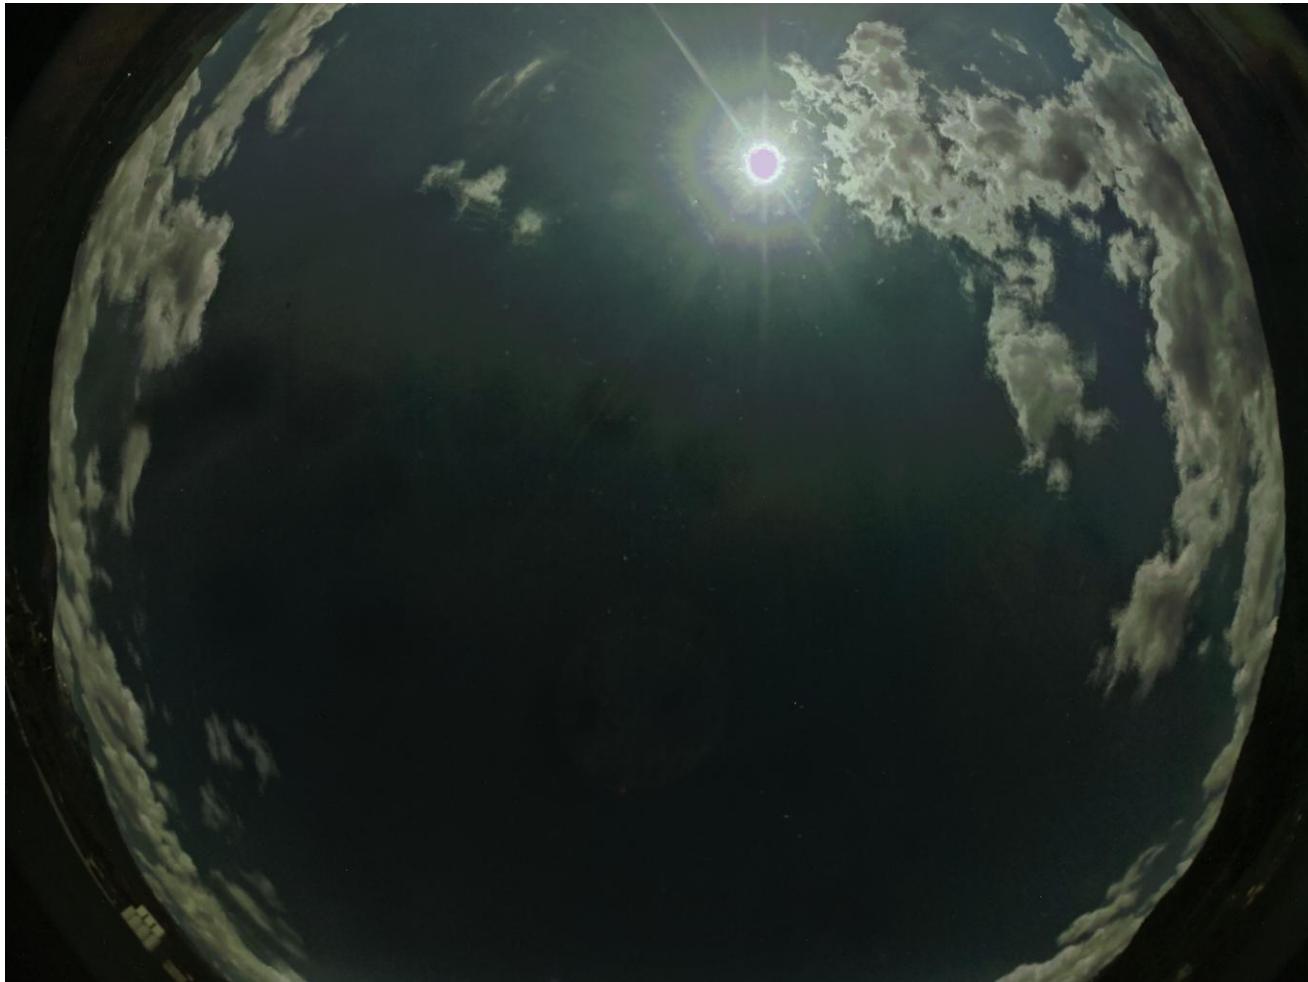
# Labelling

- **Goal**
  - Assign each sky image to the correct **cloud-cover class** (Clear, Cloudy, Overcast)
- Images were labelled manually using the same labelling criteria.
- Out of 14'000 images, 3000 images were labelled → 1000 images per class.

# Clear, cloudy or overcast?



# Clear, cloudy or overcast?



Cloudy

# Clear, cloudy or overcast?



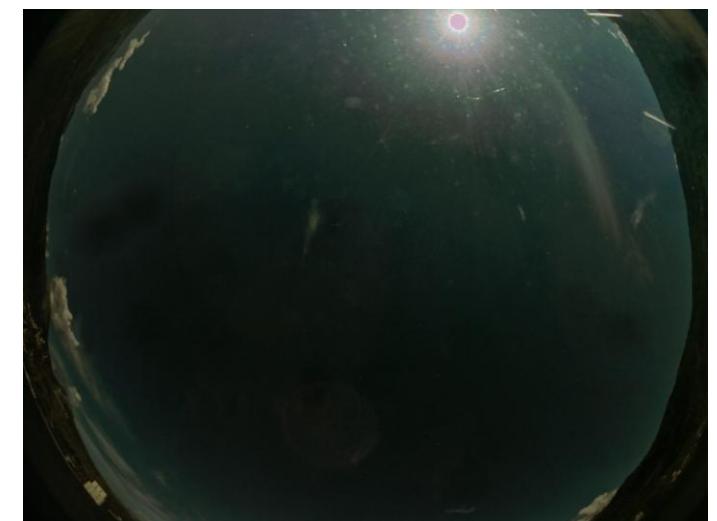
# Clear, cloudy or overcast?



Cloudy

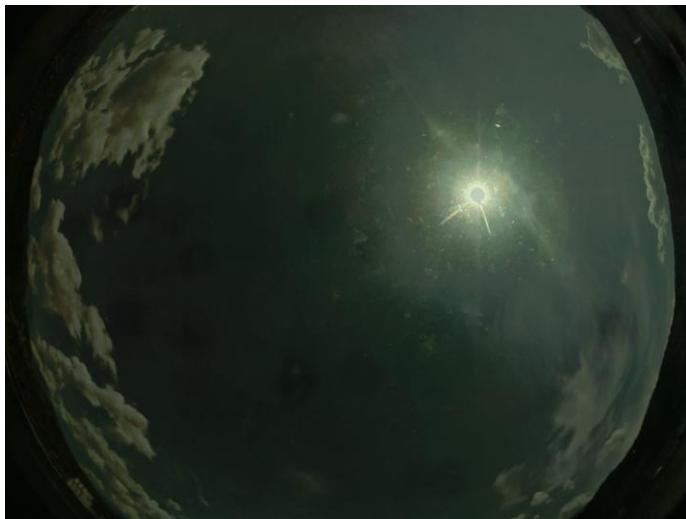
# Class Clear

- Sky is almost entirely free of clouds. Direct sunlight reaches the ground with little scattering.
- Sun is visible and unobstructed – no clouds near the sun.
- 0 – 30% cloud coverage



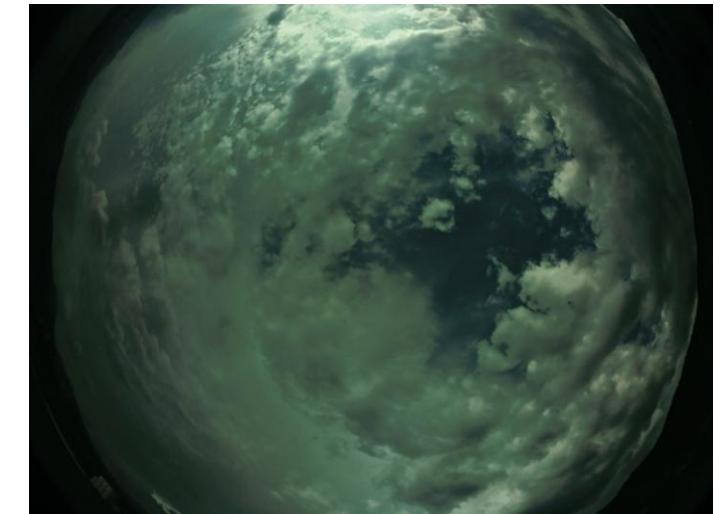
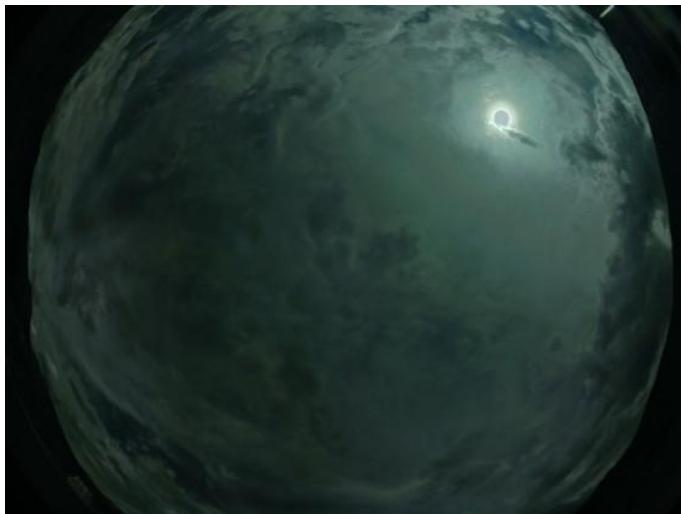
# Class Cloudy

- Sky has clouds but still shows clear patches. Sunlight sometimes gets through but is partially obstructed.
- Mixture of sky and clouds - clouds near the sun
- 30 – 70% cloud coverage



# Class Overcast

- The sun is fully covered by clouds with no visible breaks around. Cloud layer blocks direct sunlight.
- No direct sunlight
- >70% cloud coverage



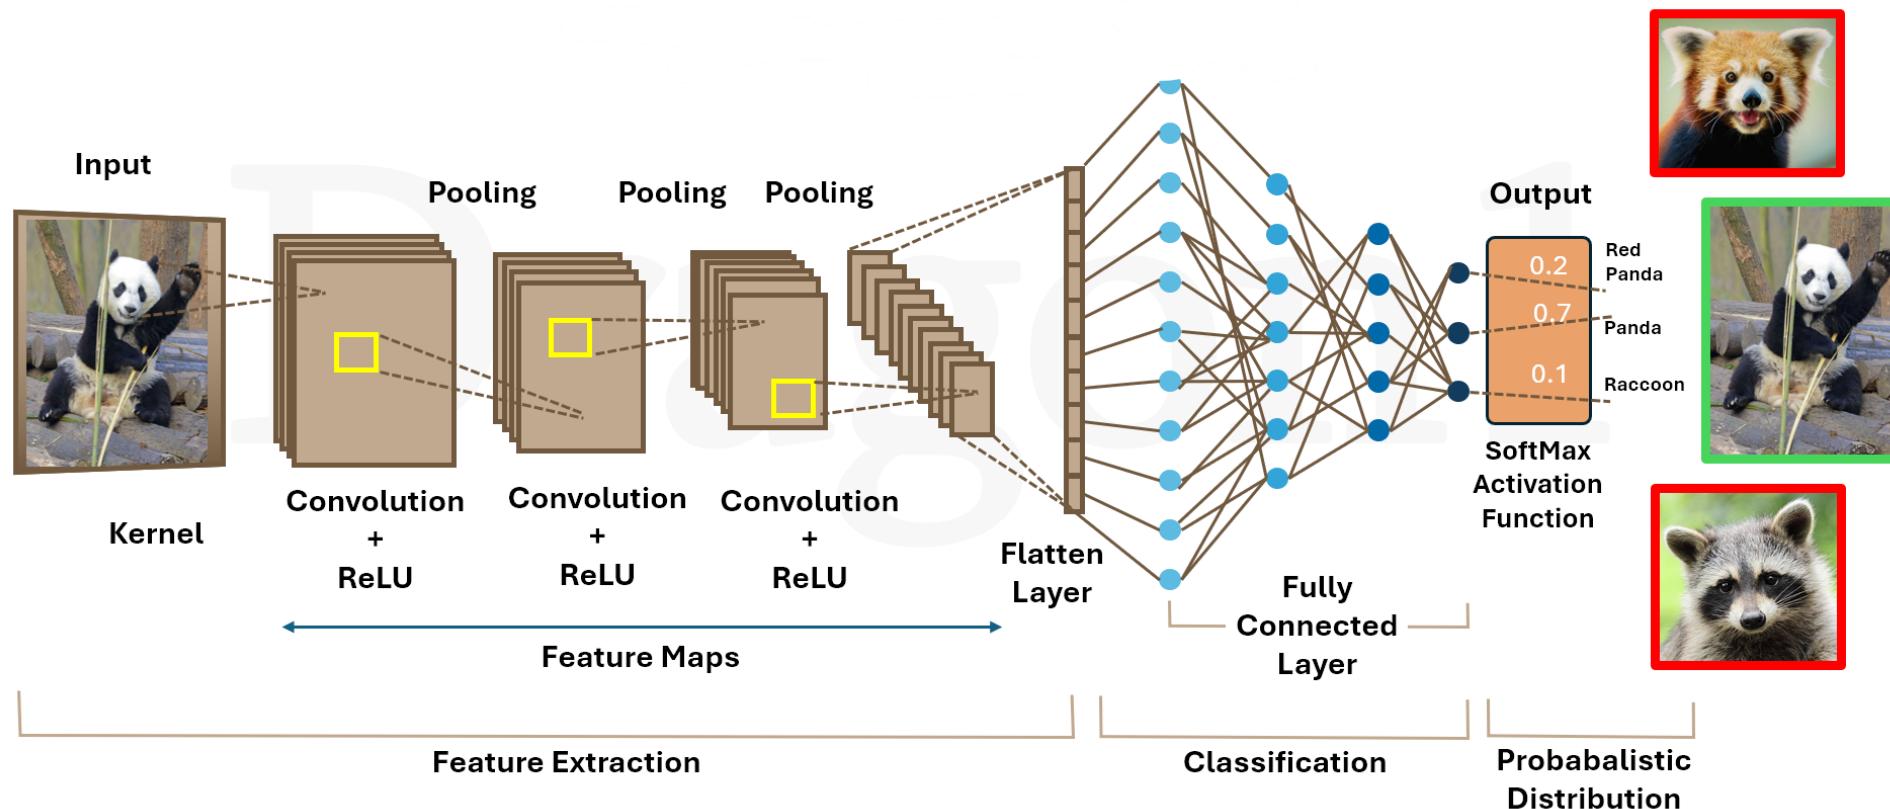
# Definitions

# Definitions

- CNN
- Transfer Learning
- Fine Tuning
- Literature Research → Model

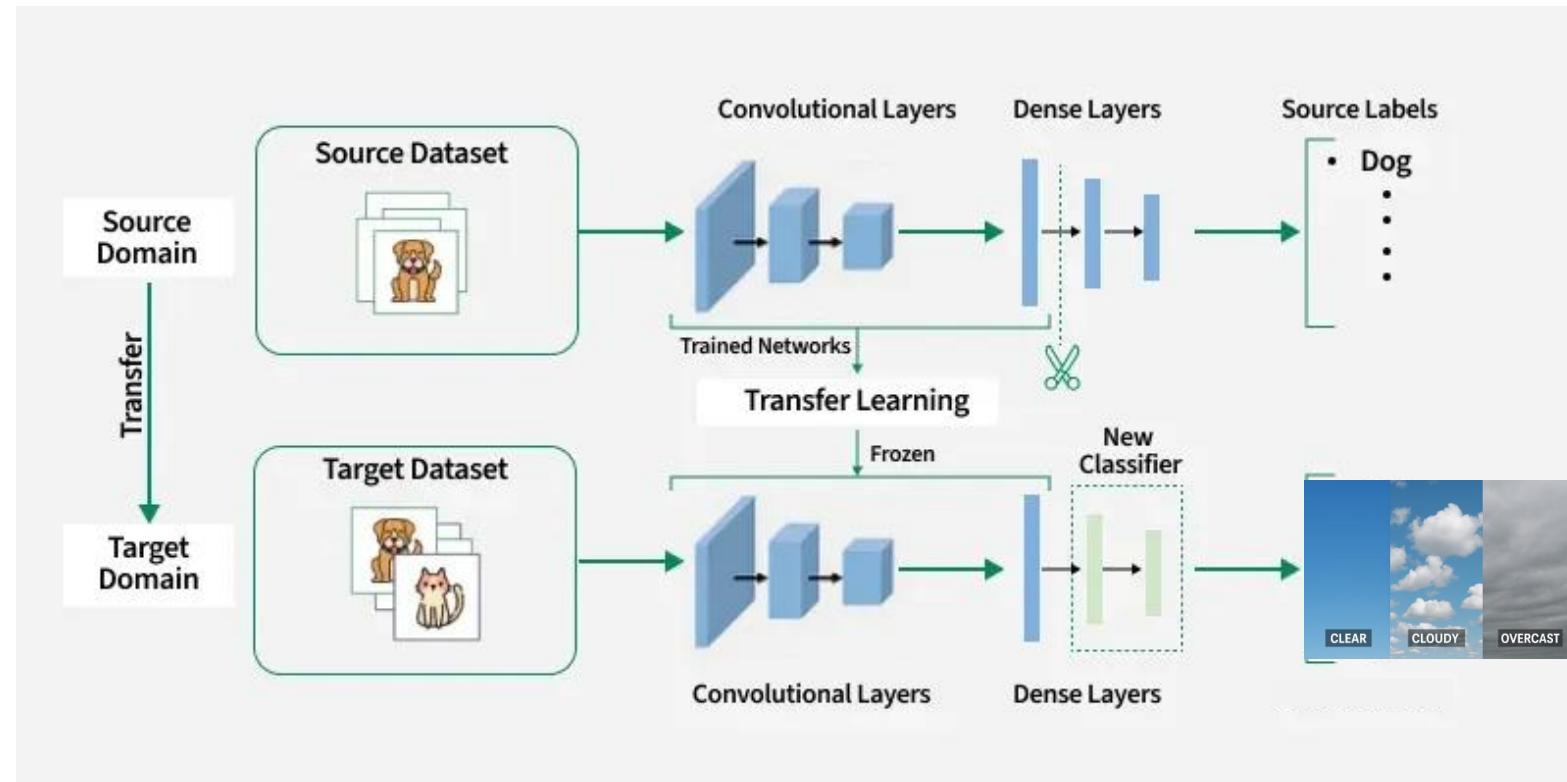
# Convolutional Neural Network (CNN)

- A CNN learns by repeatedly predicting on images, measuring how wrong it is, and adjusting its filters through backpropagation until the predictions become accurate.”



# Transfer Learning

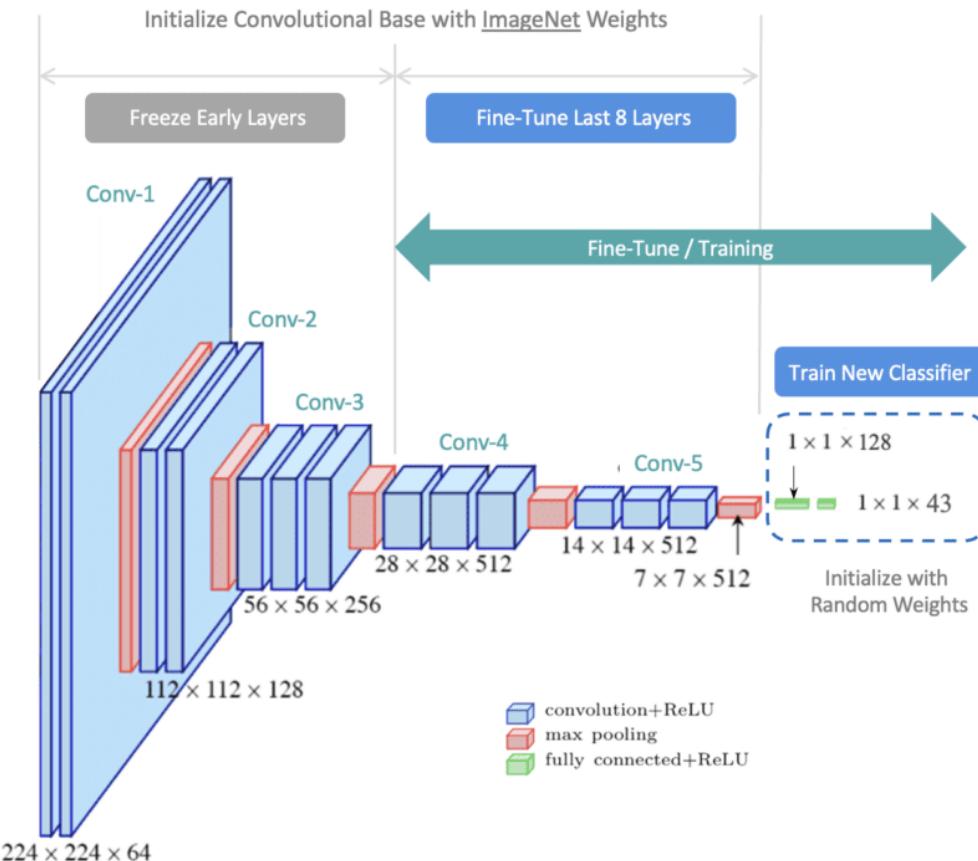
- Start with CNN trained on a large dataset  
(e.g., ImageNet with 14M images and 1000 classes)
- Already knows how to detect edges, shapes, patterns like sky, tree, objects
- Cut last classification layers, keep rest **frozen**.
- add a small **new classifier** for our own categories (clear / cloudy / overcast).



# Fine Tuning

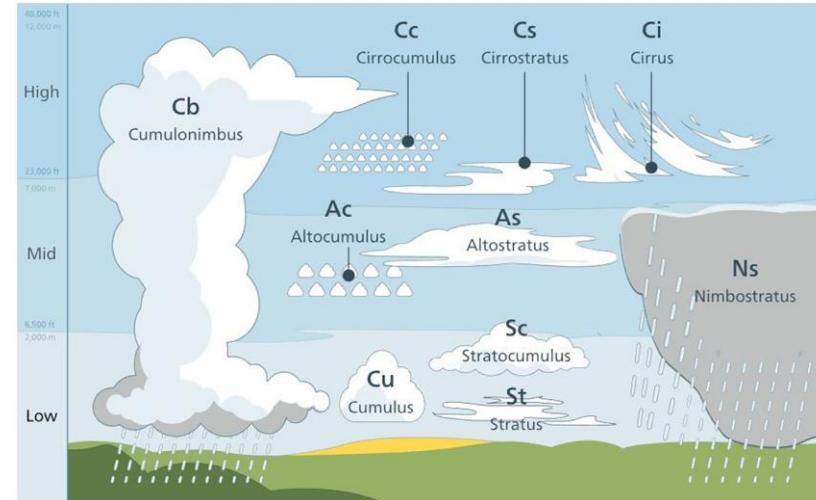
## Unlock the Power of Fine-Tuning Pre-Trained Models

- Step 1 — Start with a Pre-Trained Model (ImageNet)
- Step 2 — Freeze the Early Layers → Keep low level features
- Step 3 — Unfreeze top M layers → to make it more task-specific
  - Layers adapt to the new dataset of clouds
- Step 4 — Replace the Classifier and train

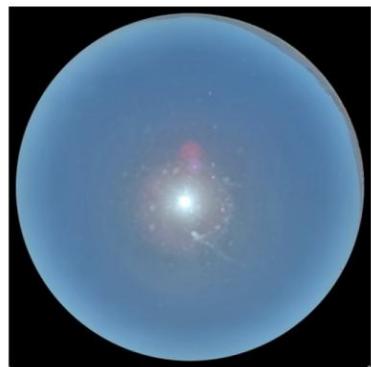


# Literature Research

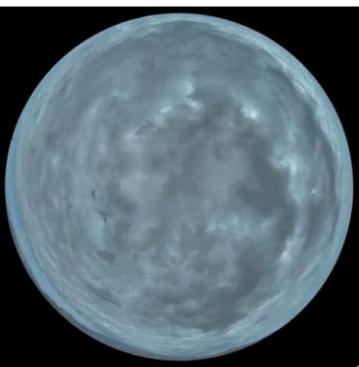
- Cloud Type Classification – Guzel et al. (2024)
  - 11 cloud types
  - Seven pretrained CNNs compared
  - Public Dataset: Cirrus Cumulus Stratus Nimbus
  - **Xception performed best with 97.66% accuracy**
- Weather Image Classification – Naufal et al. (2021)
  - Out of 4 CNN architectures tested, **Xception performed best with 90.21% accuracy**.



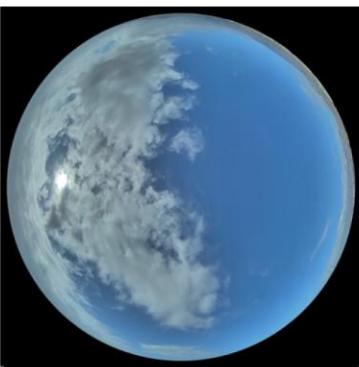
# Literature Research



(a) Clear sky class.



(b) Cloudy sky class.



(c) Partly cloudy sky class.

- Sky Image Classification – Hernández-López et al. (2024)
  - Clear / Partly Cloudy / Cloudy
- Tested EfficientNet V2 (B1, B2) and ResNet models
- Best performance:  
**EfficientNetV2-B1/B2 with 98.09% accuracy.**

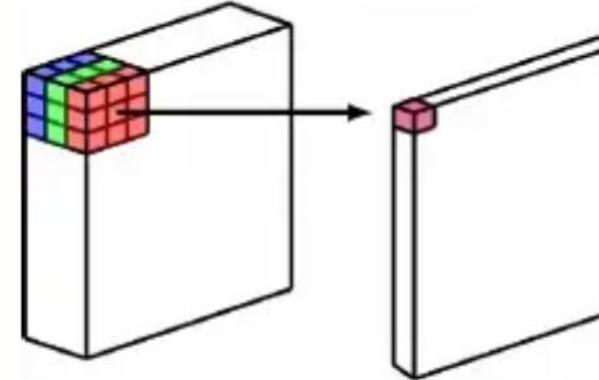
# Conclusion from Literature Research

- Initially planned to compare multiple architectures, but this exceeded the scope of the project.
  - Selected **Xception**.
- **Research Question**
  - How effectively can a CNN based on the Xception architecture classify images into **clear, cloudy and overcast** using transfer learning?*
- **Hypothesis**
  - Given that Xception achieved **97.66% accuracy** in cloud-type classification, it is expected to perform equally well — or better — on the simpler task of cloud-cover classification.

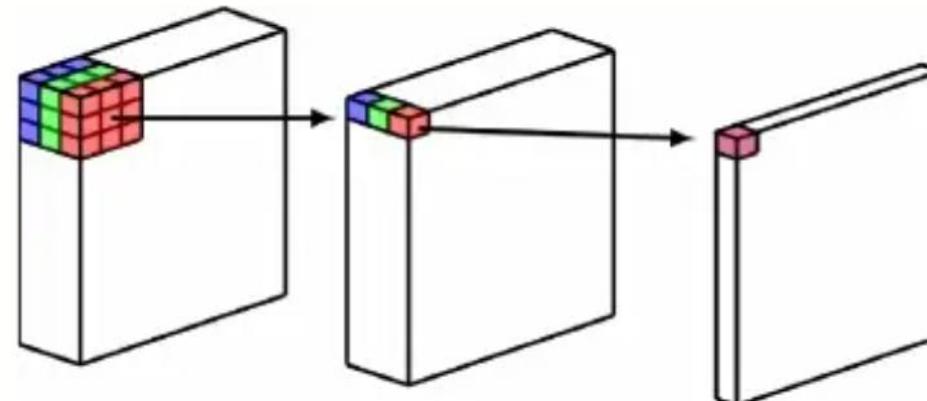
# Base Model Xception

# Xception

- Developed by François Chollet
  - 2017 at Google
- Uses **Depthwise Separable Convolutions**
  - First: each colour channel filtered separately (fine details)
  - Then: pointwise  $1 \times 1$  convolution combines information
  - Computationally less expensive
- Pre-trained on ImageNet → strong transfer learning



(a) Conventional Convolutional Neural Network



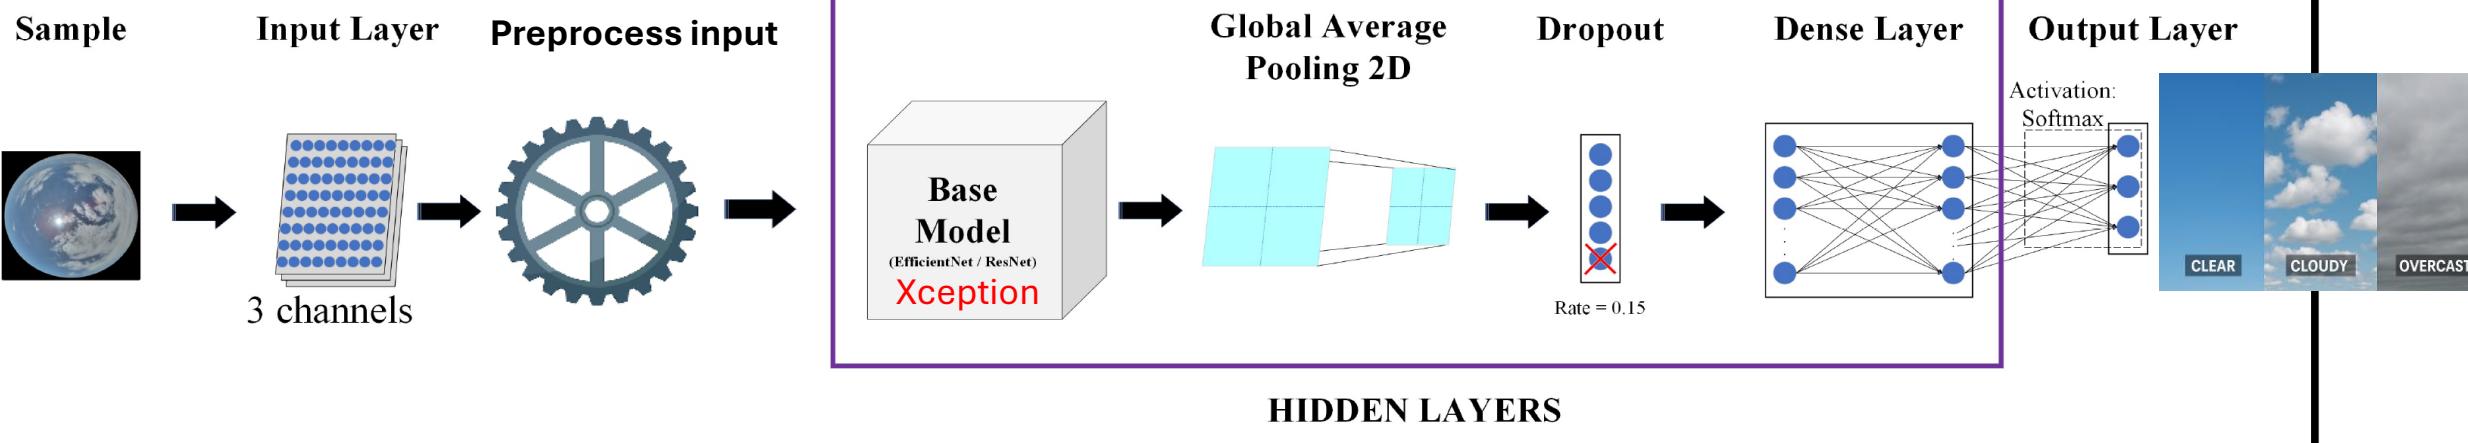
Depthwise Convolution  
Pointwise Convolution  
(b) Depthwise Separable Convolutional Neural Network

# Implementation Architecture

```
from tensorflow.keras.applications
import Xception

base = Xception(
    include_top = False,
    weights = "imagenet",
    input_shape=IMG_SIZE + (3,)
)
base.trainable = False
```

```
inputs = keras.Input(shape = IMG_SIZE + (3,))
x = preprocess_input(x)
x = base(x)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.2)(x)
outputs = layers.Dense(len(CLASS_NAMES),
activation="softmax")(x)
```



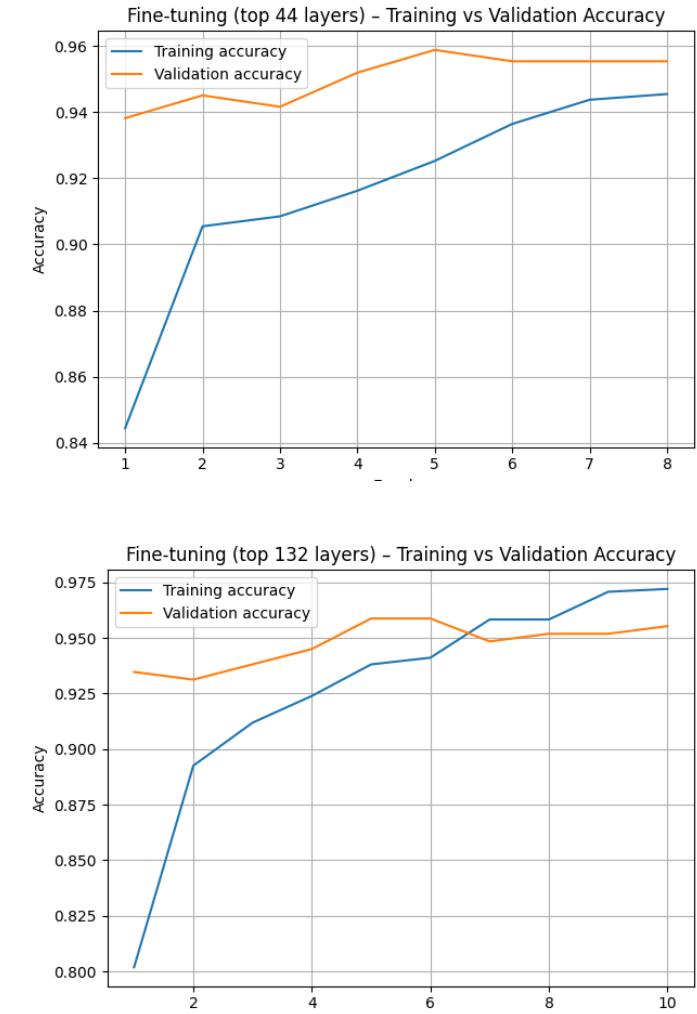
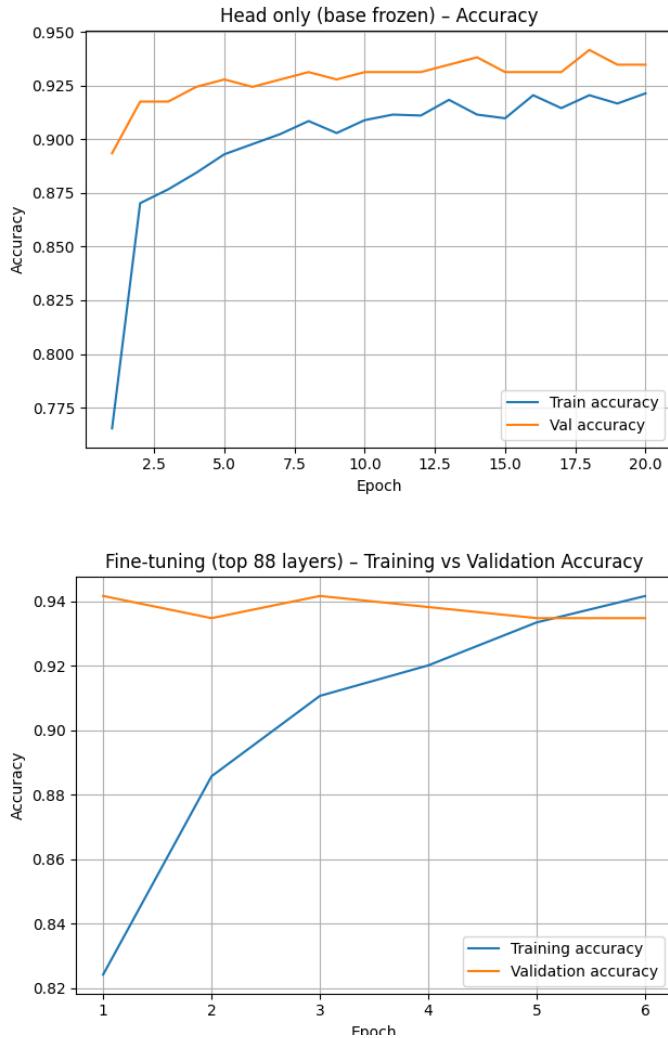
## Training and Results

# Training Setup

- 3 classes: **clear, cloudy, overcast**
- Split: 80% train / 10% validation / 10% test over 3000 labelled images.
- Metrics used:
  - **Accuracy** – how many images are correctly classified.
  - **Loss** – Sparse categorical cross entropy
  - **Precision / Recall / F1 (macro)** – average over the three classes, so each class counts equally.
- Two stages (20 Epochs):
  - **Stage 1:** train only the new head (Xception frozen)
  - **Stage 2: fine-tuning** with different numbers of unfrozen layers (44, 88, 132)

# Accuracy

- Head-only training already gives strong performance
- Early stopping triggers quickly during fine-tuning
- In FT-132 and FT-88, training accuracy grows much faster than validation accuracy and even surpasses it → Overfitting
- Validation accuracy staying stable while training accuracy rises



# Training set metrics

Train set	Accuracy	Loss	Precision	Recall	F1
<b>Base model</b>	0.9208	0.2117	0.9219	0.9206	0.9210
<b>FT - 44</b>	0.9432	0.1500	0.9436	0.9437	0.9436
<b>FT - 88</b>	0.9367	0.1667	0.9366	0.9370	0.9366
<b>FT - 132</b>	0.9712	0.0891	0.9710	0.9717	0.9712

FT – 44 = Xception Fine-tuned with 44 layers unfrozen

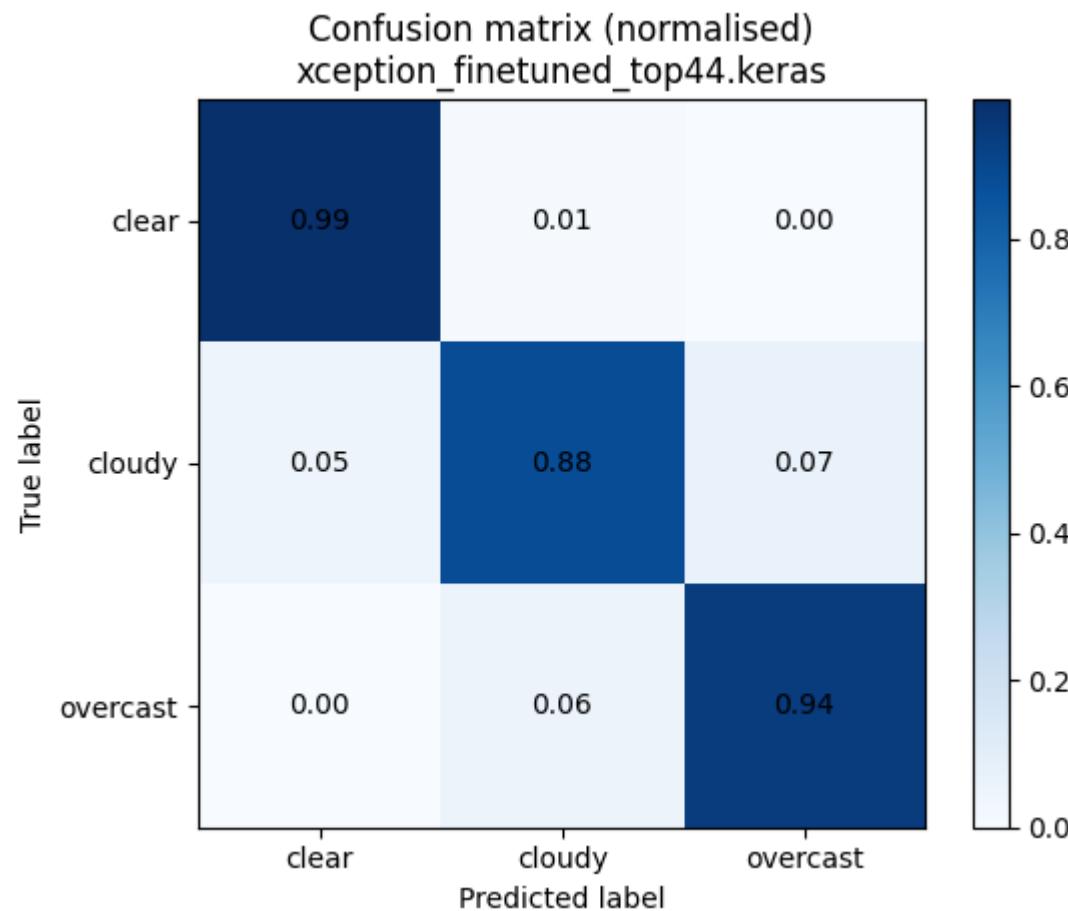
# Validation set metrics

Val set	Accuracy	Loss	Precision	Recall	F1
<b>Base model</b>	0.9367	0.1650	0.9416	0.9347	0.9367
<b>FT - 44</b>	<b>0.9588</b>	0.1358	<b>0.9626</b>	<b>0.9586</b>	<b>0.9599</b>
<b>FT - 88</b>	0.9416	0.1792	0.9429	0.9437	0.9428
<b>FT - 132</b>	0.9450	<b>0.1348</b>	0.9484	0.9450	0.9462

# Test set metrics

Test set	Accuracy	Loss	Precision	Recall	F1
Base model	0.9212	0.1981	0.9205	0.9203	0.9082
FT - 44	0.9384	0.1727	0.9374	0.9382	0.9374
FT - 88	0.9110	0.2386	0.9128	0.9128	0.9083
FT - 132	0.9315	0.1545	0.9303	0.9301	0.9230

# Test set metrics

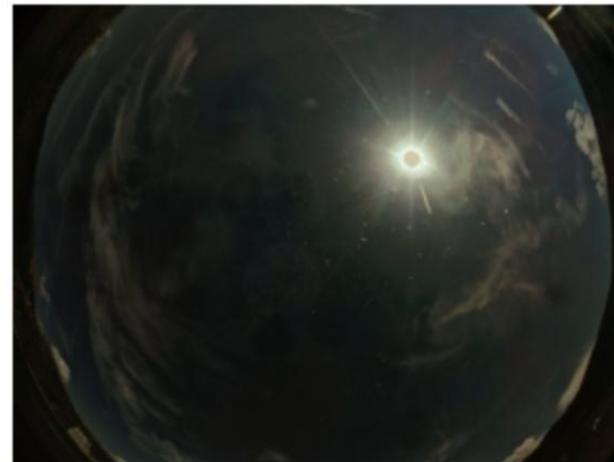


# Some misclassified images

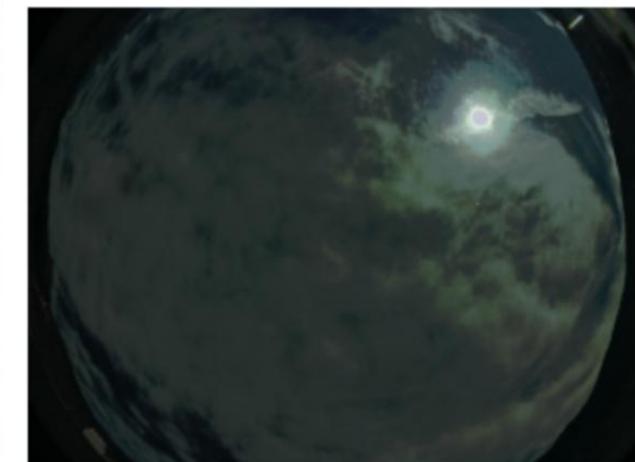
True: clear  
Pred: cloudy



True: cloudy  
Pred: clear



True: cloudy  
Pred: overcast



# Analysis of Results

- Strong classifier with all four model variants achieving high values in all metrics.
  - 91 – 97% train accuracy
  - 93 – 96% validation accuracy
  - 91 – 94 % test accuracy
  - Precision/Recall/F1 between 90% and 97%
- ✓ The model is consistently above 90% on unseen data.
- ✓ It correctly classifies clear / cloudy / overcast images with high reliability.

# Conclusion

- **Research Question**

*How effectively can a CNN based on the Xception architecture classify images into **clear, cloudy and overcast** using transfer learning?*

- **Hypothesis**

Given that Xception achieved 97.66% accuracy in cloud-type classification, it is expected to perform equally well — or better — on the simpler task of cloud-cover classification.

→ **Hypothesis Outcome:**

Although the model did **not reach the expected 97.66%**, it still achieved a **high accuracy of 93.86%**, confirming that Xception is highly effective for cloud-cover classification.

# Conclusion

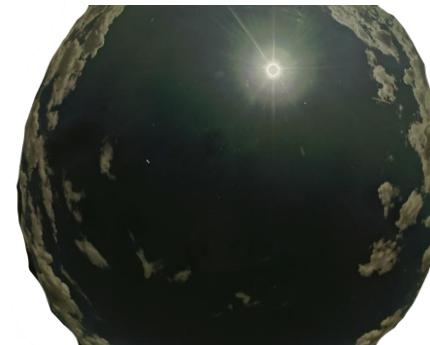
- Fine-tuning improves performance
    - **FT-44** achieves **best** result with **93.84% accuracy**.
    - **FT-88 & FT-132** show signs of **overfitting**
- Fine-tuning increases performance by a few percentage points (92% in base model), but this comes at the cost of **higher computational load, longer training times**, and the risk of **overfitting**. Whether the additional accuracy is worth the extra complexity depends entirely on the application requirements.

## Future Work

# Possible Data Preparation Steps

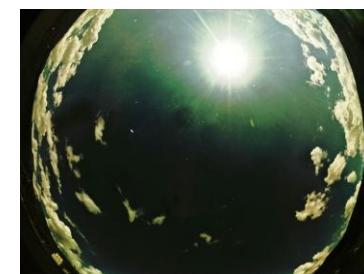
- **Apply circular mask and resize to 299 x 299**

- Remove black rim
- Cut away objects at the horizon
- Squeeze it to 299 x 299



- **Augmentation**

- Rotation
- Color shifts
- Brightness adjustments



# Future Research

- Cloudy sky → cloud type classification
- Xception very heavyweight – try base model with less parameters.  
Code is built very modularly, so can be easily replaced. E.g.  
Efficientnet V2

Model	Size (MB)	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth	Time (ms) per inference step (CPU)	Time (ms) per inference step (GPU)
Xception	88	79.0%	94.5%	22.9M	81	109.4	8.1
EfficientNetV2B2	42	80.5%	95.1%	10.2M	-	-	-
MobileNet	16	70.4%	89.5%	4.3M	55	22.6	3.4

- CNN + other model for cloud movement and irradiance prediction

*u*<sup>b</sup>

---

*b*  
UNIVERSITÄT  
BERN

That's it!  
Thank you for your attention!

# Questions

# Questions

- Do you think it was worth it to fine-tune?
- Do you have other ideas on how to predict solar irradiance? Hybrid models?
- What is the most realistic next step with the conclusions of this project?
- Do you have suggestions of other models for transfer learning?
- If you were to collect the dataset again and preprocess what would you change to make the model more robust?

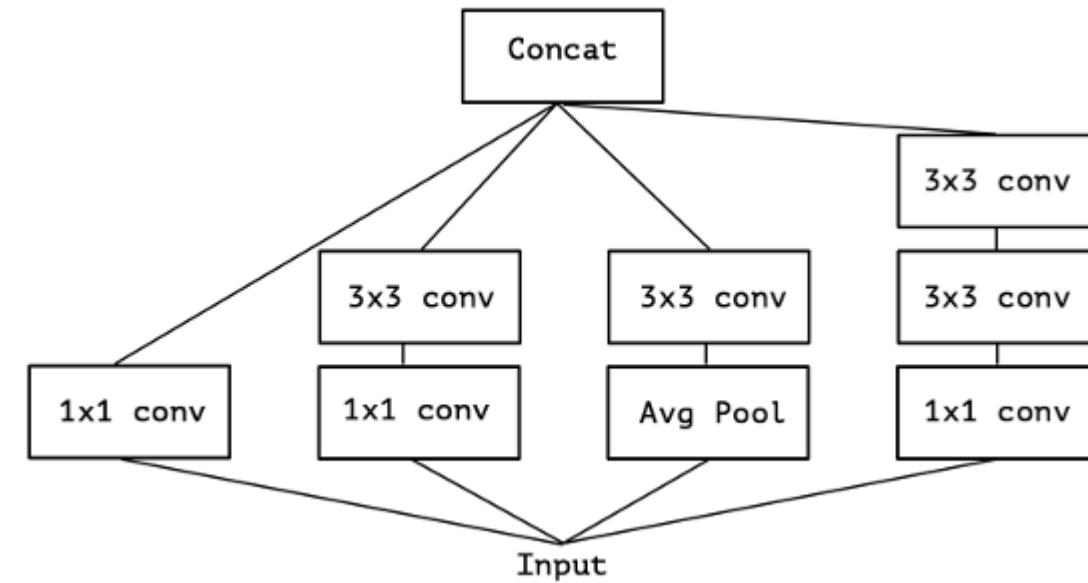
# References

- F. Chollet, "Xception: Deep Learning with Depthwise Separable Convolutions," 2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR), Honolulu, HI, USA, 2017, pp. 1800-1807, doi: 10.1109/CVPR.2017.195.
- Guzel, M., Kalkan, M., Bostancı, E., Acici, K., & Asuroglu, T. (2024). Cloud type classification using deep learning with cloud images. *PeerJ. Computer science*, 10, e1779.  
<https://doi.org/10.7717/peerj-cs.1779>
- Hernández-López, R., Travieso-González, C. M., & Ajali-Hernández, N. I. (2024). Sky Image Classification Based on Transfer Learning Approaches. *Sensors*, 24(12), 3726.  
<https://doi.org/10.3390/s24123726>
- Kalkan M, Bostancı GE, Güzel MS, Kalkan B, Özsarı Ş, Soysal Ö, Köse G. 2022. Cloudy/clear weather classification using deep learning techniques with cloud images. *Computers and Electrical Engineering* 102(7804):108271 DOI 10.1016/j.compeleceng.2022.108271.
- Naufal, M. F., & Kusuma, S. F. (2022, April). Weather image classification using convolutional neural network with transfer learning. In *AIP Conference Proceedings* (Vol. 2470, No. 1, p. 050004). AIP Publishing LLC.

# Evaluation Metrics

Metric	Definition
Accuracy	Percentage of all predictions that are correct.
Loss	Measures how wrong the model is during training (lower is better).
Precision	Of all images predicted as a class, how many are actually correct.
Recall	Of all real images of a class, how many the model correctly finds.
F1-Score	Harmonic mean of precision and recall; high only if both are high.

# 132 length vs 81 depth



# Why SVM for Raindrop

- Because raindrop detection is a **small, highly imbalanced, and very visual problem**.  
An SVM works better when you have:
- **few training images**
- **only two classes**
- **high-quality feature embeddings** (like MobileNetV2's 1280-D vectors)
- a **simple linear boundary** between classes
- Softmax requires more data and tends to overfit quickly.  
SVMs, especially with pre-trained embeddings, are **very strong for small binary problems**.

# My journey

- Problem description
  - Why do we need it
- Preprocessing
- Transfer Learning
  - Xception architecture (warum Xception?)
- My CNN
  - First time 87%
  - Extract wrong images and see who was right me or CNN
  - # Part 1: train only the small "head" (fast, safe).#
  - Part 2: fine-tune the top layers of Xception (better accuracy).#
  - Then: evaluate (accuracy, precision, recall, F1), confusion matrix
  - Fine tuning with different layers frozen
- How to make it better
  - Preprocessing steps
- Question round

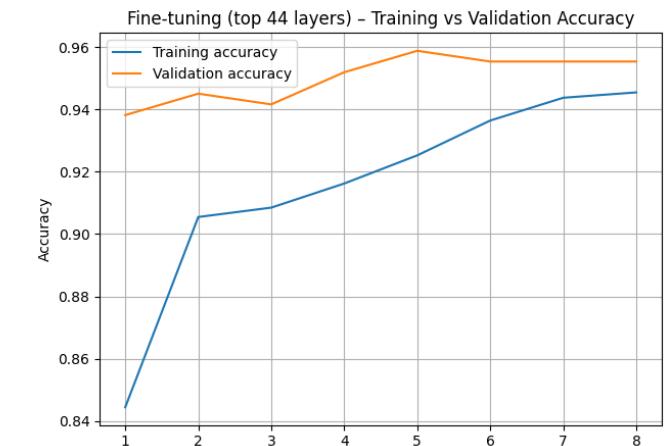
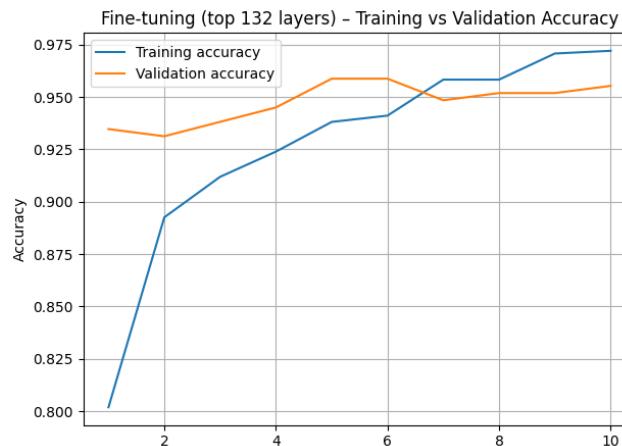
# Fine Tuning – How it Works

- Freezing

# FINE TUNING

- The fine tuning process took over 1 hour
- So I split it and saved my progress in a model keras file after training the head and after fine tuning.
- Problem is fine tuning causes overfitting.
- When putting images from the new server, it might be as accurate

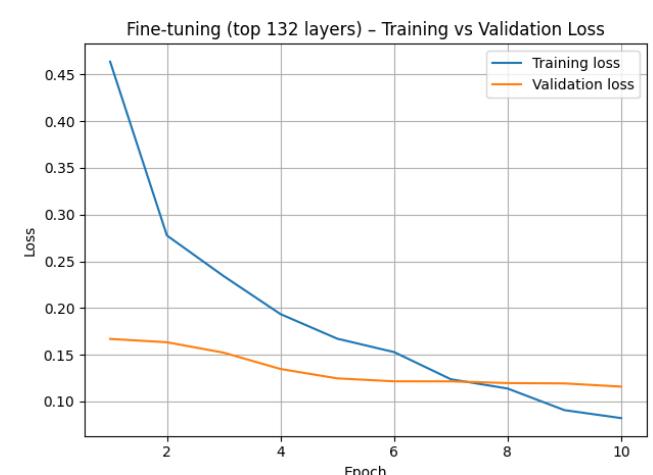
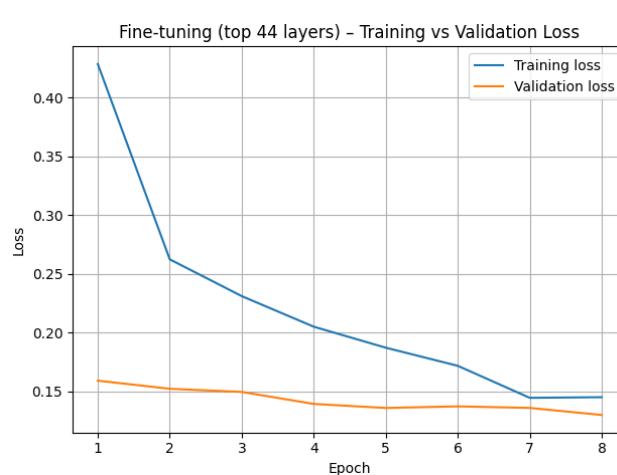
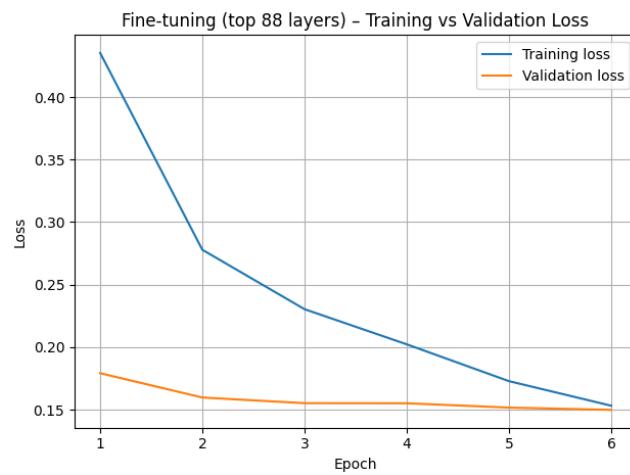
# Accuracy



# ETH Study

- <https://www.research-collection.ethz.ch/server/api/core/bitstreams/47bd7a0d-5c7f-41b6-8752-7bceb3407f35/content>

# Loss



*u*<sup>b</sup>

---

*b*  
UNIVERSITÄT  
BERN

# In a Broader Context